


SPRINT DEMO

Plugging your domain agent into the platform

Two endpoints. Three levels. One black box.

- 1 The HTTP seam
- 2 Evaluating skill selection

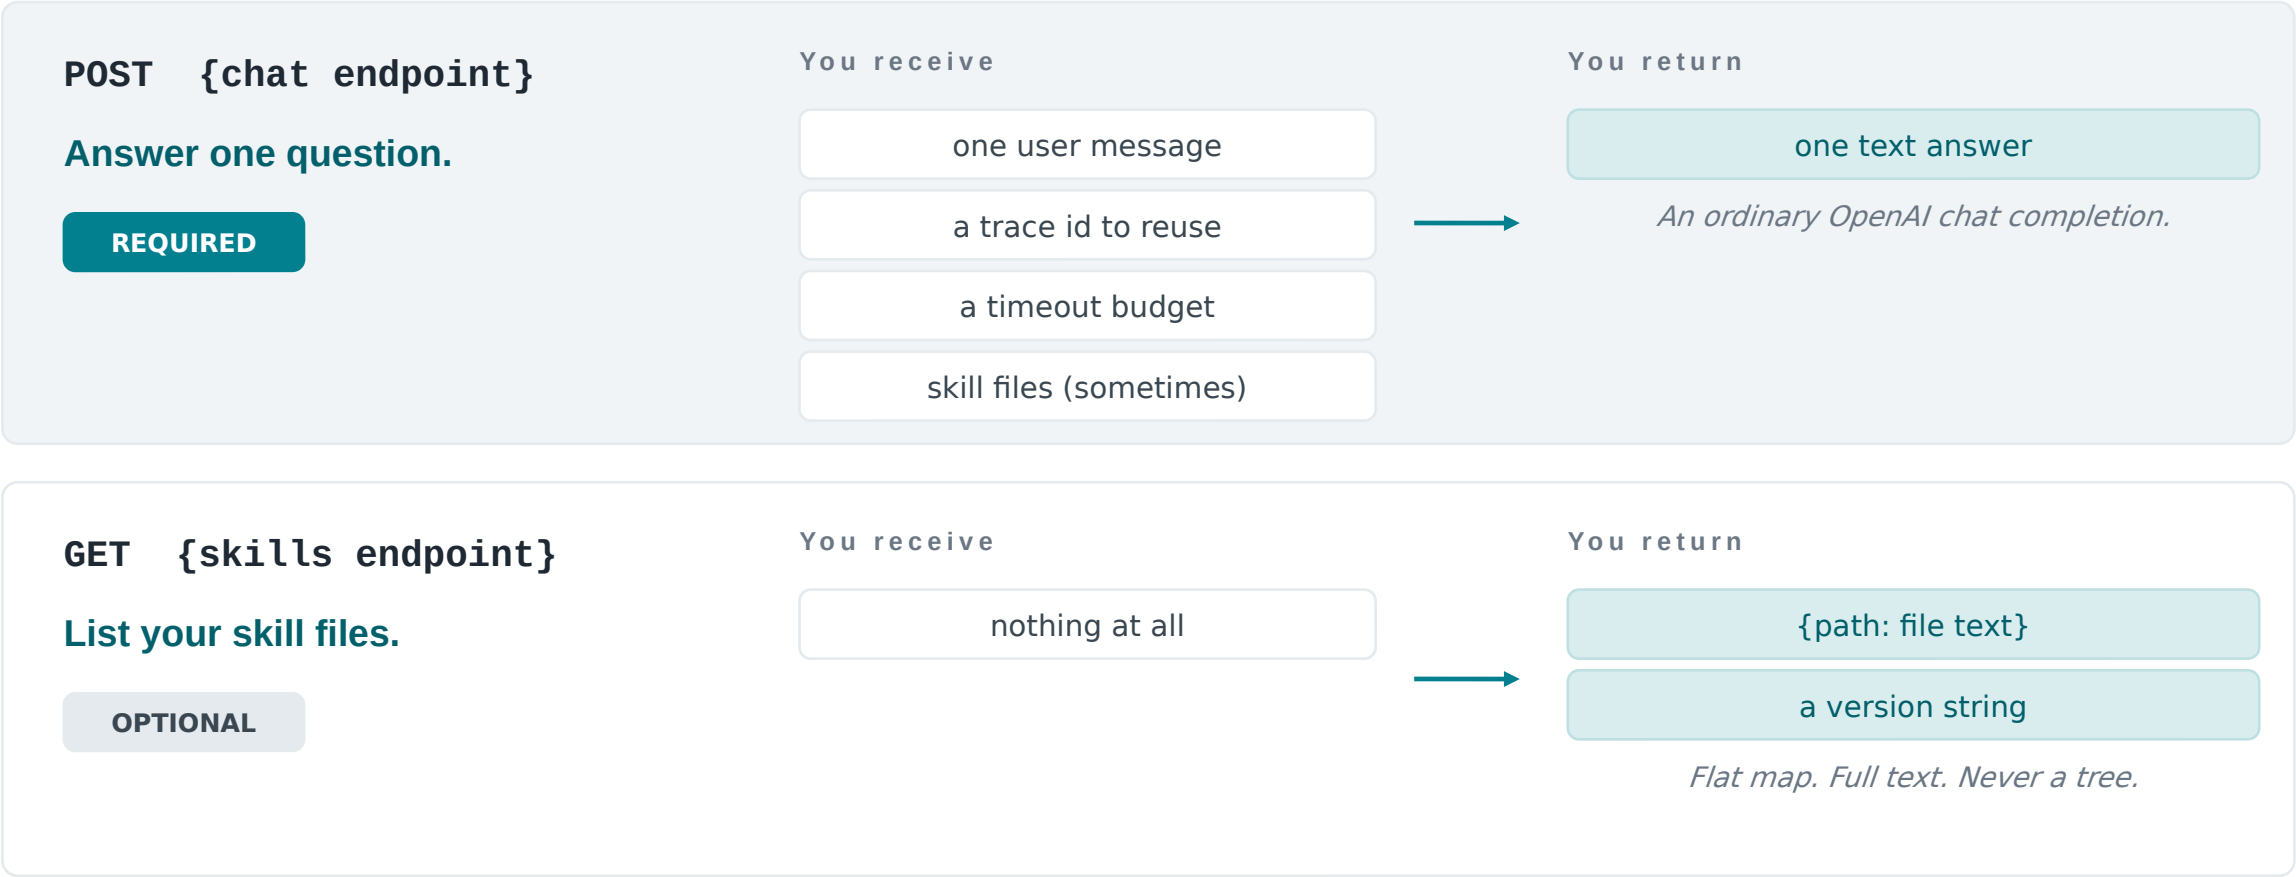
Where your agent meets the platform



Left of the seam is ours. Right of it is yours. The seam itself is what this sprint made general.

Two endpoints

Two absolute URLs. Only the first is required.



No path is imposed: different prefixes, gateways, even different hosts are all fine.

Inside one chat call

```
{
  "model": "default",
  "messages": [
    {"role": "user", "content": "What was ACME's
      outstanding balance at the end of Q2?"}
  ],
  "stream": false,
  "skill_studio": {
    "timeout_s": 115.0,
    "trace_data": {"trace_id": "9f3e11c8...", ...},
    "skills": {"billing/SKILL.md": "..."}
  }
}
```

The trace id is the load-bearing one.

What your side must do

- 1 Reuse the trace id**
not your own
- 2 Honour timeout_s**
then 504
- 3 Apply skills, if sent**
those files only
- 4 Answer with text**
never empty

The response is an ordinary chat completion. Empty content is a failed question, not a wrong answer — return a 5xx with a reason instead.

The skills you run with — and the ones we lend you

```
{
  "version": "a1b2c3d",
  "skills": {
    "billing/SKILL.md": "...",
    "billing/references/refunds.md": "...",
    "reporting/SKILL.md": "..."
  }
}
```

Flat map — one key per file

Full text — never truncated

{ } is an answer; 5xx is a failure

version moves when behaviour moves

Send that same map back on a chat call and it becomes an override — three states:

no skills key

SKILL.md

refs

...

Your own files, as normal.

"skills": { ... }

SKILL.md

refs

These files only, for this call.

"skills": { }

nothing

No skill files at all.

Replacement, never a patch. Never persisted. Test for the key, not for truthiness — { } is falsy.

You do not have to build all of it

The platform works with what you give it, and names what is missing.

1

A chat endpoint

what most agents already have

You get

Evaluation

Score an eval set. Read every verdict.

2

+ skills, + override

one route, one behaviour

You get

The playground

Edit skills. Ask one question. Catch drift.

3

+ our trace id, in Langfuse

not a third endpoint

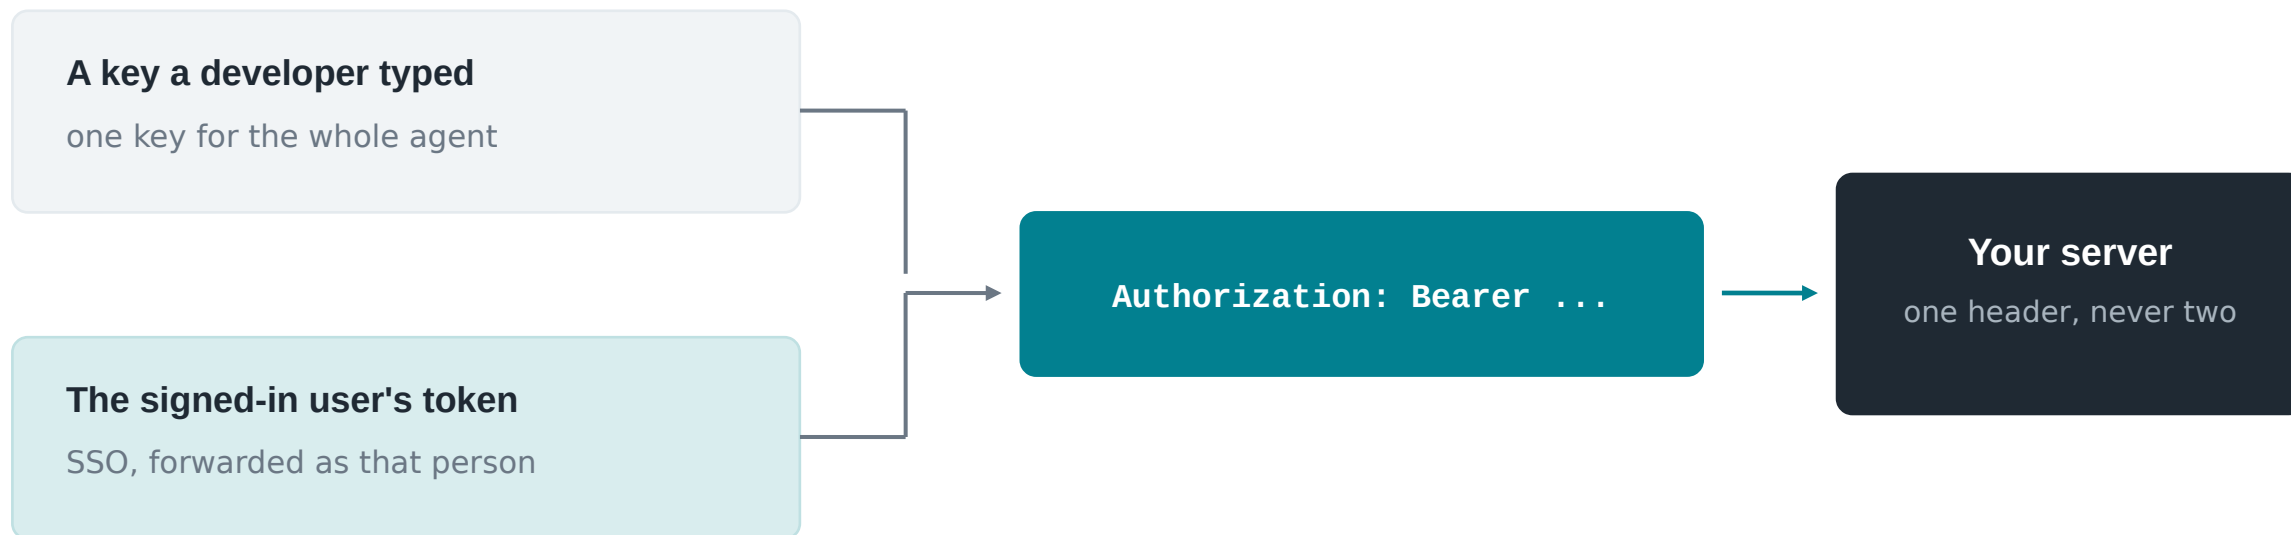
You get

Optimization

Hundreds of scored rollouts against candidates.

If your server needs to know who is asking

Authentication unlocks no feature — but the platform can send a credential.



Why the second one matters

It is the only mode that tells you who is asking, so an agent with its own per-user authorization keeps it.

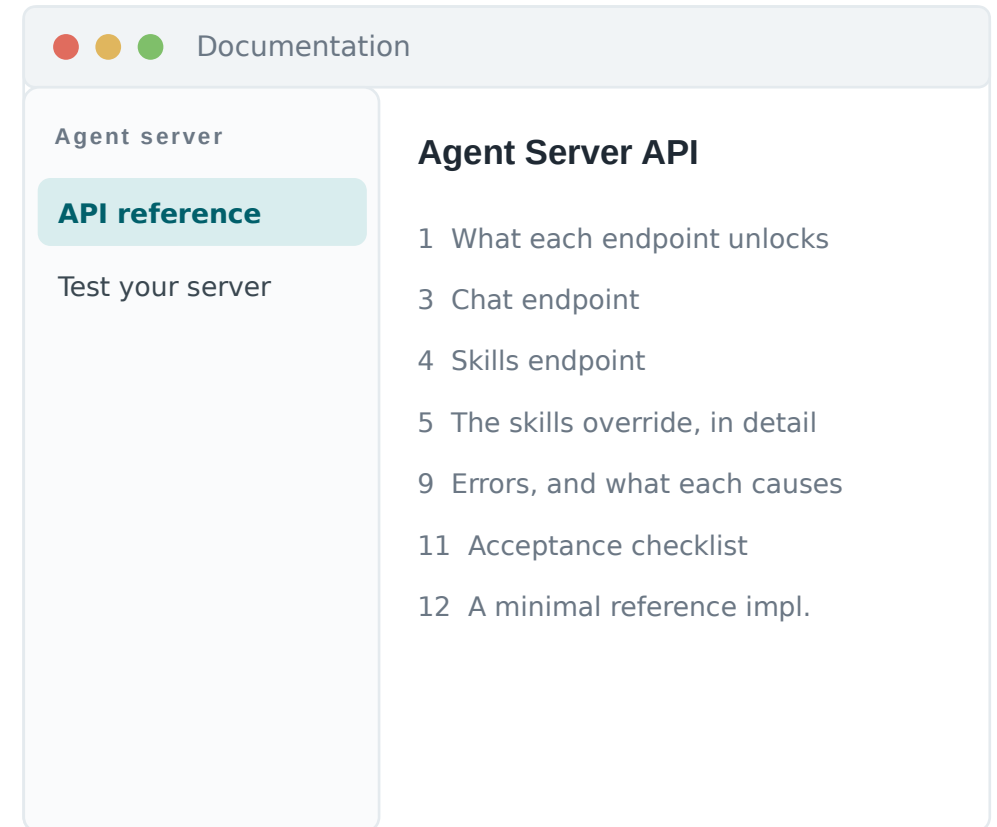
Two things to check

Same realm and same audience. Tokens are re-minted mid-run — cache on the claims, not the token text.

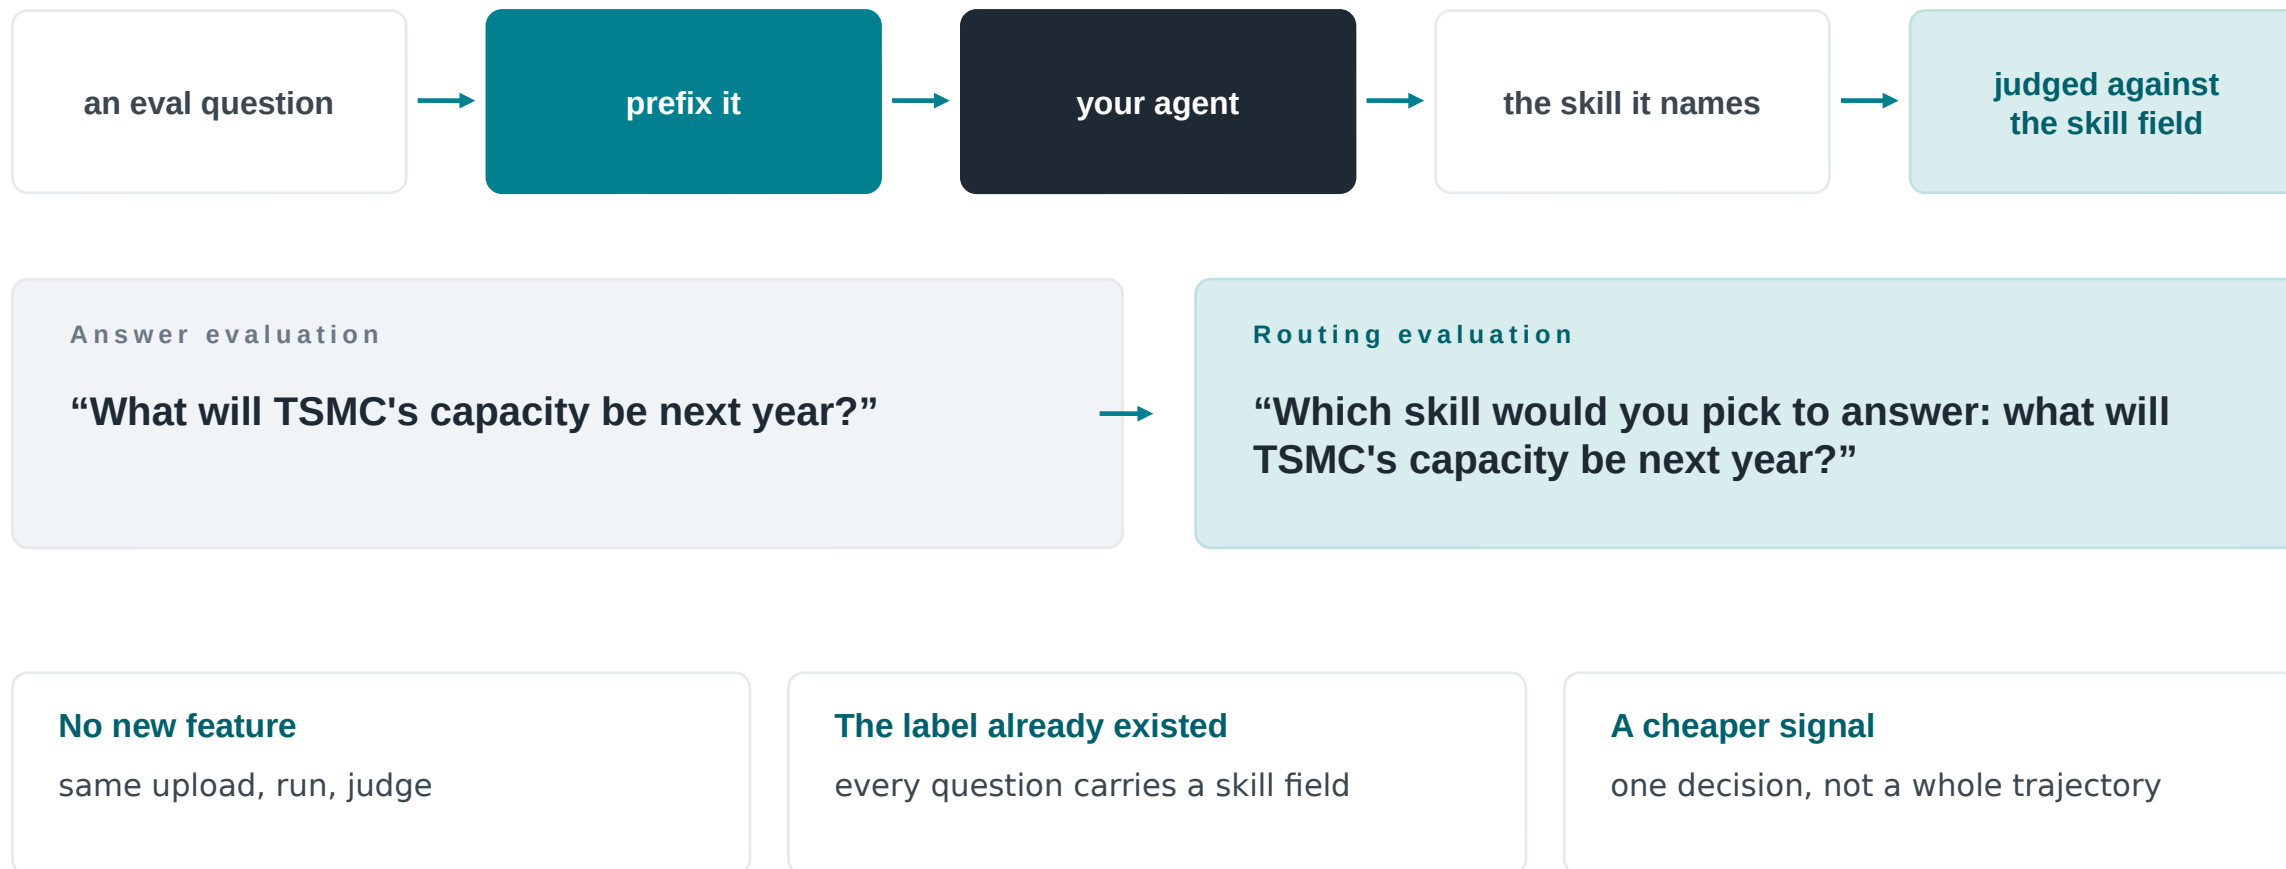
You should probably not write this by hand

- 1 Open the API reference**
one self-contained page
- 2 Hand it to your coding assistant**
it writes the server
- 3 Run Test your server**
it names the level you reached

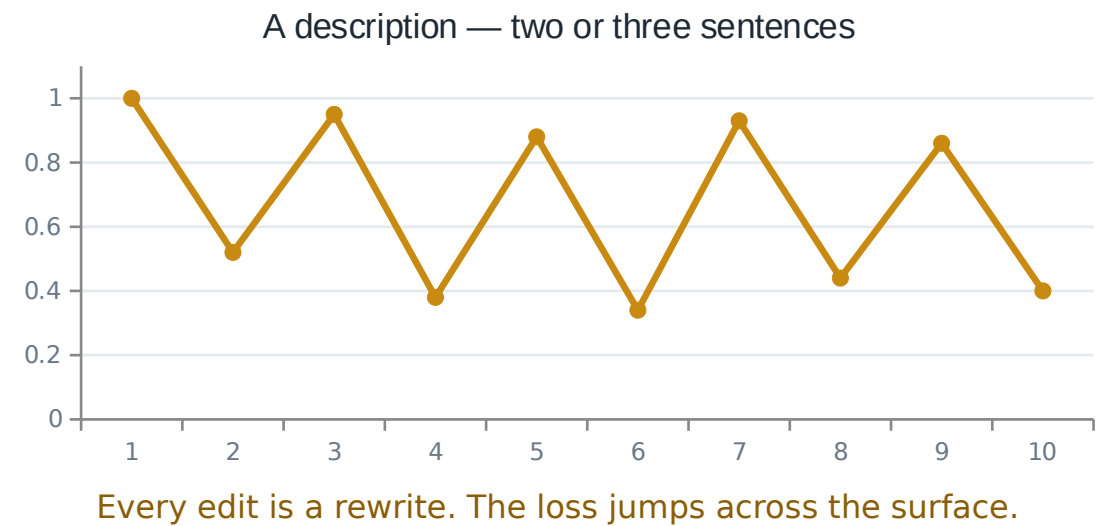
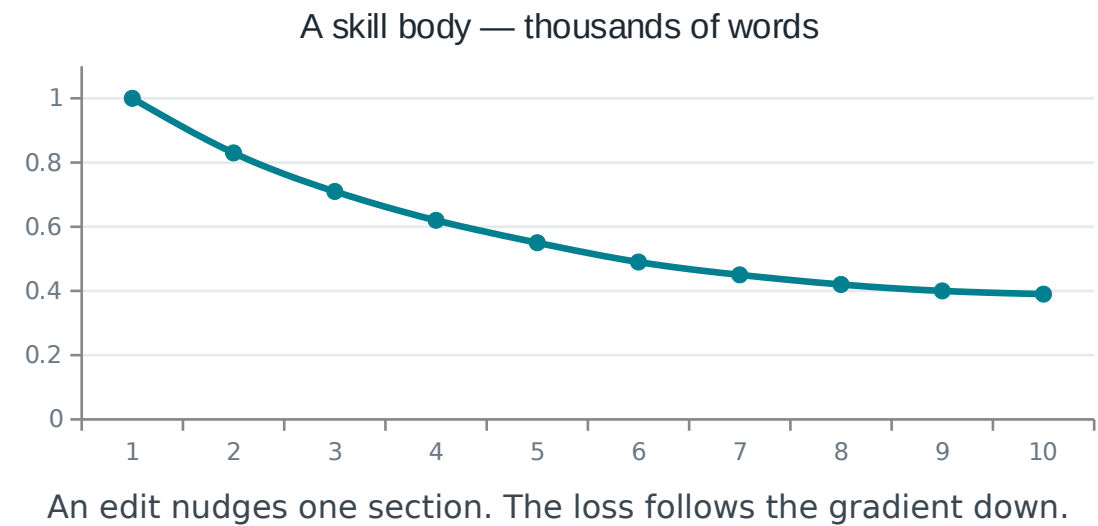
→ **DEMO**



The same eval loop, pointed at skill selection



Why the SkillOpt recipe did not transfer



! Still in flight

results next sprint

What I'd like from you

Start at level 1.

1

Bring the chat endpoint you have

evaluation needs nothing else

2

Hand the API reference to your assistant

it was written for that

3

Prove it with Test your server

it names the level you reached

4

Tell me where the doc made you guess

that is a bug in the doc

Questions?